

Assignment4_9628_jzeiders

November 11, 2024

1 John Zeiders (jzeiders) - Assignment 4

```
[ ]: import numpy as np
      from numpy.linalg import inv, det
      import pandas as pd
```

2 Part 1

```
[ ]: def Estep(data, prob, mean, Sigma):
      """
      E-step: Calculate responsibilities

      Parameters:
      -----
      data: ndarray, shape (n, p)
          The input data matrix
      prob: ndarray, shape (G,)
          Mixing proportions
      mean: ndarray, shape (p, G)
          Mean vectors for each component
      Sigma: ndarray, shape (p, p)
          Shared covariance matrix

      Returns:
      -----
      responsibilities: ndarray, shape (n, G)
          Matrix of posterior probabilities  $P(Z_i=k|x_i)$ 
      """
      n, p = data.shape
      G = prob.shape[0]

      # Initialize matrix to store densities
      densities = np.zeros((n, G))

      # Precompute constants related to the covariance matrix
      try:
          inv_Sigma = np.linalg.inv(Sigma) # Inverse of covariance matrix
```

```

        det_Sigma = np.linalg.det(Sigma) # Determinant of covariance matrix
        if det_Sigma <= 0:
            raise np.linalg.LinAlgError("Covariance matrix is not positive_
↪definite.")
        except np.linalg.LinAlgError as e:
            raise ValueError("Invalid covariance matrix. " + str(e))

        # Compute the normalization constant for the multivariate normal_
↪distribution
        norm_const = 1.0 / (np.power((2 * np.pi), p / 2) * np.sqrt(det_Sigma))

        # Compute densities for each Gaussian component
        for k in range(G):
            # Extract the mean vector for the k-th component
            mu_k = mean[:, k] # Shape: (p,)

            # Compute the difference between each data point and the mean vector
            diff = data - mu_k # Shape: (n, p)

            # Compute the exponent term for the multivariate normal PDF
            # This represents: -0.5 * (x - mu_k)^T * inv_Sigma * (x - mu_k)
            # Efficiently computed using einsum
            exponent = -0.5 * np.einsum('ij,jk,ik->i', diff, inv_Sigma, diff) #_
↪Shape: (n,)

            # Compute the density for the k-th component
            densities[:, k] = norm_const * np.exp(exponent) # Shape: (n,)

        # Multiply by mixing proportions
        weighted_densities = densities * prob[np.newaxis, :]

        # Normalize to get responsibilities
        responsibilities = weighted_densities / np.sum(weighted_densities, axis=1)[:
↪, np.newaxis]

        return responsibilities

def Mstep(data, responsibilities):
    """
    M-step: Update parameters

    Parameters:
    -----
    data: ndarray, shape (n, p)
        The input data matrix
    responsibilities: ndarray, shape (n, G)
        Matrix of posterior probabilities

```

```

Returns:
-----
prob: ndarray, shape (G,)
    Updated mixing proportions
mean: ndarray, shape (p, G)
    Updated mean vectors
Sigma: ndarray, shape (p, p)
    Updated shared covariance matrix
"""
n, p = data.shape
G = responsibilities.shape[1]

# Update mixing proportions
prob = np.mean(responsibilities, axis=0)

# Update means
mean = np.zeros((p, G))
for k in range(G):
    mean[:, k] = np.sum(responsibilities[:, k:k+1] * data, axis=0) / np.
↪sum(responsibilities[:, k])

# Update shared covariance matrix
Sigma = np.zeros((p, p))
for k in range(G):
    diff = data - mean[:, k] # Shape: (n, p)
    weighted_diff = responsibilities[:, k].reshape(n, 1) * diff # Shape: ↪
↪(n, p)
    Sigma += weighted_diff.T @ diff # Shape: (p, p)
Sigma /= n

return prob, mean, Sigma

def loglik(data, prob, mean, Sigma):
    """
    Compute log-likelihood

    Parameters:
    -----
    data: ndarray, shape (n, p)
        The input data matrix
    prob: ndarray, shape (G,)
        Mixing proportions
    mean: ndarray, shape (p, G)
        Mean vectors
    Sigma: ndarray, shape (p, p)
        Shared covariance matrix

```

```

Returns:
-----
ll: float
    Log-likelihood value
"""
n = data.shape[0]
G = prob.shape[0]

# Initialize array for component densities
densities = np.zeros((n, G))
d = data.shape[1]
inv_Sigma = np.linalg.inv(Sigma)
det_Sigma = np.linalg.det(Sigma)
norm_const = 1.0 / (np.power((2 * np.pi), d / 2) * np.sqrt(det_Sigma))

# Compute densities for each component
for k in range(G):
    # Compute the difference between data and the k-th mean
    diff = data - mean[:, k] # Shape: (n_samples, d)

    # Compute the exponent term:  $-0.5 * (diff @ inv\_Sigma * diff)$ .
    ↪sum(axis=1)
    # Efficient computation using einsum for element-wise multiplication ↪
    ↪and summation
    exponent = -0.5 * (diff @ inv_Sigma * diff).sum(axis=1)

    # Compute the density for the k-th component
    densities[:, k] = prob[k] * norm_const * np.exp(exponent)

# Sum over components and take log
ll = np.sum(np.log(np.sum(densities, axis=1)))

return ll

def myEM(data, G, initial_params, itmax):
    """
    Main EM algorithm function

    Parameters:
    -----
    data: ndarray, shape (n, p)
        The input data matrix
    G: int
        Number of components
    initial_params: dict
        Dictionary containing initial parameters:

```

```

        'prob': mixing proportions
        'mean': mean vectors
        'Sigma': covariance matrix
    itmax: int
        Maximum number of iterations

    Returns:
    -----
    prob: ndarray, shape (G,)
        Final mixing proportions
    mean: ndarray, shape (p, G)
        Final mean vectors
    Sigma: ndarray, shape (p, p)
        Final shared covariance matrix
    ll: float
        Final log-likelihood value
    """
    # Initialize parameters
    prob = initial_params['prob']
    mean = initial_params['mean']
    Sigma = initial_params['Sigma']

    for _ in range(itmax):
        # E-step
        responsibilities = Estep(data, prob, mean, Sigma)

        # M-step
        prob, mean, Sigma = Mstep(data, responsibilities)

    # Compute final log-likelihood
    ll = loglik(data, prob, mean, Sigma)

    return prob, mean, Sigma, ll

```

```

[ ]: def test_G2(data):
    """Test EM algorithm with G=2 components"""
    n, p = data.shape

    # Set initial parameters as specified
    p1 = 10/n
    prob = np.array([p1, 1-p1])

    mean = np.zeros((p, 2))
    mean[:, 0] = np.mean(data[:10], axis=0) # mean of first 10 samples
    mean[:, 1] = np.mean(data[10:], axis=0) # mean of remaining samples

    # Calculate initial Sigma

```

```

diff1 = data[:10] - mean[:, 0]
diff2 = data[10:] - mean[:, 1]
Sigma = (diff1.T @ diff1 + diff2.T @ diff2) / n

initial_params = {
    'prob': prob,
    'mean': mean,
    'Sigma': Sigma
}

# Run EM algorithm
final_prob, final_mean, final_Sigma, final_ll = myEM(data, 2, □
↪initial_params, 20)

return final_prob, final_mean, final_Sigma, final_ll

def test_G3(data):
    """Test EM algorithm with G=3 components"""
    n, p = data.shape

    # Set initial parameters
    p1, p2 = 10/n, 20/n
    prob = np.array([p1, p2, 1-p1-p2])

    mean = np.zeros((p, 3))
    mean[:, 0] = np.mean(data[:10], axis=0) # mean of first 10 samples
    mean[:, 1] = np.mean(data[10:30], axis=0) # mean of next 20 samples
    mean[:, 2] = np.mean(data[30:], axis=0) # mean of remaining samples

    # Calculate initial Sigma
    diff1 = data[:10] - mean[:, 0]
    diff2 = data[10:30] - mean[:, 1]
    diff3 = data[30:] - mean[:, 2]
    Sigma = (diff1.T @ diff1 + diff2.T @ diff2 + diff3.T @ diff3) / n

    initial_params = {
        'prob': prob,
        'mean': mean,
        'Sigma': Sigma
    }

    # Run EM algorithm
    final_prob, final_mean, final_Sigma, final_ll = myEM(data, 3, □
↪initial_params, 20)

    return final_prob, final_mean, final_Sigma, final_ll

```

```
[ ]: # Load faithful dataset
file_path = 'https://liangfgithub.github.io/Data/faithful.dat' # Update this_
    ↪ path as necessary
data = np.loadtxt(file_path, skiprows=1, dtype=float, usecols=(1,2))
print("Data Shape:", data.shape)
print("First 5 Data Points:\n", data[:5])
```

Data Shape: (272, 2)

First 5 Data Points:

```
[[ 3.6  79. ]
 [ 1.8  54. ]
 [ 3.333 74. ]
 [ 2.283 62. ]
 [ 4.533 85. ]]
```

```
[ ]: # Test G=2
print("\nTesting with G=2:")
prob2, mean2, Sigma2, ll2 = test_G2(data)

print("\nprob")
print(prob2)
print("\nmean")
print(mean2)
print("\nSigma")
print(Sigma2)
print("\nloglik")
print(ll2)
```

Testing with G=2:

prob

```
[0.04297883 0.95702117]
```

mean

```
[[ 3.49564188  3.48743016]
 [76.79789154 70.63205853]]
```

Sigma

```
[[ 1.29793612 13.92433626]
 [13.92433626 182.58009247]]
```

loglik

```
-1289.5693549424104
```

```
[ ]: # Test G=3
print("\nTesting with G=3:")
prob3, mean3, Sigma3, ll3 = test_G3(data)
```

```

print("\nprob")
print(prob3)
print("\nmean")
print(mean3)
print("\nSigma")
print(Sigma3)
print("\nloglik")
print(ll3)

```

Testing with G=3:

```

prob
[0.04363422 0.07718656 0.87917922]

mean
[[ 3.51006918  2.81616674  3.54564083]
 [77.10563811 63.35752634 71.25084801]]

Sigma
[[ 1.26015772 13.51153756]
 [13.51153756 177.96419105]]

loglik
-1289.350958862739

```

3 Part 2

```

[ ]: def forward_pass(data, w, A, B):
    """
    Compute forward probabilities  $\alpha(t,j)$ 

    Args:
        data: T-by-1 observation sequence (1D array)
        w: Initial state distribution (mz,)
        A: Transition matrix (mz-by-mz)
        B: Emission matrix (mz-by-mx)

    Returns:
        alpha: Forward probabilities (T-by-mz)
    """
    T = len(data)
    mz = len(w)
    alpha = np.zeros((T, mz))

    # Initialize first time step

```



```

alpha[0, :] = w * B[:, data[0]]

# Forward recursion
for t in range(1, T):
    alpha[t, :] = np.dot(alpha[t - 1, :], A) * B[:, data[t]].T

return alpha

def backward_pass(data, A, B):
    """
    Compute backward probabilities beta(t,j)

    Args:
        data: T-by-1 observation sequence (1D array)
        A: Transition matrix (mz-by-mz)
        B: Emission matrix (mz-by-mx)

    Returns:
        beta: Backward probabilities (T-by-mz)
    """
    T = len(data)
    mz = A.shape[0]
    beta = np.zeros((T, mz))

    # Initialize last time step
    beta[-1, :] = 1

    # Backward recursion
    for t in range(T - 2, -1, -1):
        beta[t, :] = np.dot(A, (B[:, data[t + 1]] * beta[t + 1, :]))

    return beta

def BW_onestep(data, w, A, B):
    """
    One step of the Baum-Welch algorithm (E-step + M-step)

    Args:
        data: T-by-1 observation sequence (1D array)
        w: Initial state distribution (mz,)
        A: Current transition matrix (mz-by-mz)
        B: Current emission matrix (mz-by-mx)

    Returns:
        A_new: Updated transition matrix
        B_new: Updated emission matrix
    """

```

```

T = len(data)
mz = A.shape[0]
mx = B.shape[1]

# E-step: Compute forward and backward probabilities
alpha = forward_pass(data, w, A, B)
beta = backward_pass(data, A, B)

# Compute  $\xi(t, i, j) = P(z_t=i, z_{t+1}=j \mid x_{1:T})$ 
xi = np.zeros((T - 1, mz, mz))
for t in range(T - 1):
    denominator = np.dot(np.dot(alpha[t, :], A) * B[:, data[t + 1]], beta[t + 1, :])
    numerator = alpha[t, :, np.newaxis] * A * B[:, data[t + 1]] * beta[t + 1, :]
    xi[t, :, :] = numerator / denominator

# Compute state probabilities (gamma)
gammas_j = np.zeros((T, mz))
gammas_j[:-1] = np.sum(xi, axis=2)
# Fix for last time step - should use alpha and beta
gammas_j[-1] = (alpha[-1] * beta[-1]) / np.sum(alpha[-1] * beta[-1])

# M-step: Update parameters
# Update A
A_new = np.zeros_like(A)
for i in range(mz):
    for j in range(mz):
        numerator = np.sum(xi[:, i, j])
        denominator = np.sum(xi[:, i, :])
        A_new[i, j] = numerator / denominator

# Update B (vectorized version)
B_new = np.zeros_like(B)
for l in range(mx):
    mask = (data == l)
    B_new[:, l] = np.sum(gammas_j[mask], axis=0) / (np.sum(gammas_j, axis=0) + 1e-300)

# Verify row stochasticity
assert np.allclose(np.sum(A_new, axis=1), 1, rtol=1e-5), "A matrix not row stochastic"
assert np.allclose(np.sum(B_new, axis=1), 1, rtol=1e-5), "B matrix not row stochastic"

return A_new, B_new

```

```

def myBW(data, initial_params, itmax):
    """
    Main Baum-Welch algorithm with specified initial parameters

    Args:
        data: T-by-1 observation sequence (1D array)
        initial_params: Dictionary containing initial parameters:
            'w': Initial state distribution (mz,)
            'A': Initial transition matrix (mz-by-mz)
            'B': Initial emission matrix (mz-by-mx)
        itmax: Maximum number of iterations

    Returns:
        A_final: Final transition matrix
        B_final: Final emission matrix
        ll_final: Final log likelihood
    """
    # Ensure data is a 1D array
    data = np.asarray(data).flatten()

    # Extract initial parameters
    w = initial_params["w"]
    A = initial_params["A"]
    B = initial_params["B"]

    for iteration in range(itmax):
        # Compute log likelihood
        alpha = forward_pass(data, w, A, B)
        current_ll = np.log(np.sum(alpha[-1]))

        # Update parameters
        A, B = BW_onestep(data, w, A, B)

    return A, B, current_ll

def myViterbi(data, w, A, B):
    """
    Compute the most likely state sequence using the Viterbi algorithm

    Args:
        data: T-by-1 observation sequence (1D array)
        w: Initial state distribution (mz,)
        A: Transition matrix (mz-by-mz)
        B: Emission matrix (mz-by-mx)

    Returns:

```

```

    path: Most likely state sequence
    max_prob: Probability of the most likely path
"""
# Ensure data is 1D array
data = np.asarray(data).flatten()

T = len(data)          # Length of sequence
mz = len(w)            # Number of hidden states

# Initialize tables
V = np.zeros((mz, T))  # Viterbi table
bp = np.zeros((mz, T), dtype=int) # Backpointer table

# Initialize first column of Viterbi table
V[:, 0] = np.log(w) + np.log(B[:, data[0]])

# Forward pass: Fill tables
for t in range(1, T):
    for j in range(mz):
        # Calculate probabilities for all possible previous states
        probs = V[:, t-1] + np.log(A[:, j]) + np.log(B[j, data[t]])
        # Find maximum probability and its index
        V[j, t] = np.max(probs)
        bp[j, t] = np.argmax(probs)

# Backward pass: Retrieve the most likely path
path = np.zeros(T, dtype=int)
# Find the most likely final state
path[-1] = np.argmax(V[:, -1])
max_prob = V[path[-1], -1]

# Backtrack through the sequence
for t in range(T-2, -1, -1):
    path[t] = bp[path[t+1], t+1]

return path, max_prob

```

```

[ ]: # Load HMM data
data_url = 'https://liangfgithub.github.io/Data/coding4_part2_data.txt' #
    ↳ Update if necessary
data = pd.read_table(data_url, sep="\s+", header=None).values.flatten() - 1
print("HMM Data Shape:", data.shape)
print("First 5 Observations:\n", data[:5])

# Load true state sequence (for comparison)
Z_url = 'https://liangfgithub.github.io/Data/Coding4_part2_Z.txt' # Update if
    ↳ necessary

```

```

Z_true = pd.read_table(Z_url, sep="\s+", header=None).values.flatten() - 1 # Assuming states are 1-indexed
Z_true = Z_true[~np.isnan(Z_true)]
print("True State Sequence Shape:", Z_true.shape)
print("First 5 True States:\n", Z_true[:5])

# Initialize parameters
mz = 2 # number of hidden states
mx = 3 # number of observation symbols

w_initial = np.array([0.5, 0.5])
A_initial = np.array([[0.5, 0.5], [0.5, 0.5]])
B_initial = np.array([[1/9, 3/9, 5/9], [1/6, 2/6, 3/6]])

initial_params = {
    'w': w_initial,
    'A': A_initial,
    'B': B_initial
}

# Run Baum-Welch algorithm
itmax = 100
A_final, B_final, ll_final = myBW(data, initial_params, itmax)

print("\nFinal Transition Matrix A:")
print(A_final)
print("\nFinal Emission Matrix B:")
print(B_final)
print("\nFinal Log-Likelihood:")
print(ll_final)

```

HMM Data Shape: (200,)

First 5 Observations:

[1 2 2 2 2]

True State Sequence Shape: (200,)

First 5 True States:

[0. 0. 0. 0. 0.]

Final Transition Matrix A:

[[0.49793938 0.50206062]

[0.44883431 0.55116569]]

Final Emission Matrix B:

[[0.22159897 0.20266127 0.57573976]

[0.34175148 0.17866665 0.47958186]]

Final Log-Likelihood:

-202.3062728417872

```
[ ]: # Run Viterbi algorithm
path, max_prob = myViterbi(data, w_initial, A_final, B_final)

print("\nMost Likely State Sequence:")
print(path + 1)

print("\nProbability of the Most Likely Path:")
print(max_prob)

# Compare with true states
matches = (path == Z_true).sum()
total = len(Z_true)
accuracy = matches / total
print(f"\nViterbi Path Accuracy: {accuracy * 100:.2f}%")
```

Most Likely State Sequence:

```
[1 1 1 1 1 1 1 2 1 1 1 1 1 2 2 1 1 1 1 1 1 2 2 2 2 2 1 1 1 1 1 1 1 2 1 1
 1 1 1 1 1 1 2 2 1 1 1 1 1 1 2 2 2 1 1 1 1 2 2 2 2 1 1 1 1 1 1 1 1 2 2 2 2
 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 1 1 1 1 2 1 1 1 1 1 1 1 1 1 1 1 1 1 2 1
 1 1 1 2 1 1 1 1 1 1 2 2 2 2 2 2 2 2 2 2 2 1 1 1 1 1 1 1 1 2 2 2 2 2 2 2
 2 2 2 1 1 1 2 2 2 2 2 2 1 1 1 1 1 2 2 2 2 2 2 2 2 1 1 1 2 2 2 1 1 1 1 1
 1 1 1 2 2 2 2 2 1 1 1 1 1 1 1]
```

Probability of the Most Likely Path:

-318.56245145850596

Viterbi Path Accuracy: 100.00%

```
[ ]: # Initialize B uniformly
B_uniform = np.full((mz, mx), 1/3)
initial_params_uniform = {
    'w': w_initial,
    'A': A_initial,
    'B': B_uniform
}

# Run Baum-Welch for 20 iterations
print("\nRunning Baum-Welch with Uniform B for 20 Iterations:")
A_final_20, B_final_20, ll_final_20 = myBW(data, initial_params_uniform,
    ↪itmax=20)

print("\nTransition Matrix A after 20 Iterations:")
print(A_final_20)
print("\nEmission Matrix B after 20 Iterations:")
print(B_final_20)
print("\nLog-Likelihood after 20 Iterations:")
```

```

print(ll_final_20)

# Run Baum-Welch for 100 iterations
print("\nRunning Baum-Welch with Uniform B for 100 Iterations:")
A_final_100, B_final_100, ll_final_100 = myBW(data, initial_params_uniform,
    ↪itmax=100)

print("\nTransition Matrix A after 100 Iterations:")
print(A_final_100)
print("\nEmission Matrix B after 100 Iterations:")
print(B_final_100)
print("\nLog-Likelihood after 100 Iterations:")
print(ll_final_100)

B_uniform[0, :] = np.array([[1/4, 1/2, 1/4]])
B_uniform[1, :] = np.array([[1/4, 1/2, 1/4]])
initial_params_uniform = {
    'w': w_initial,
    'A': A_initial,
    'B': B_uniform
}

# Run Baum-Welch for 20 iterations
print("\nRunning Baum-Welch with Uniform B for 20 Iterations:")
A_final_20, B_final_20, ll_final_20 = myBW(data, initial_params_uniform,
    ↪itmax=20)

print("\nTransition Matrix A after 20 Iterations:")
print(A_final_20)
print("\nEmission Matrix B after 20 Iterations:")
print(B_final_20)
print("\nLog-Likelihood after 20 Iterations:")
print(ll_final_20)

# Run Baum-Welch for 100 iterations
print("\nRunning Baum-Welch with Uniform B for 100 Iterations:")
A_final_100, B_final_100, ll_final_100 = myBW(data, initial_params_uniform,
    ↪itmax=100)

print("\nTransition Matrix A after 100 Iterations:")
print(A_final_100)
print("\nEmission Matrix B after 100 Iterations:")
print(B_final_100)
print("\nLog-Likelihood after 100 Iterations:")
print(ll_final_100)

```

Running Baum-Welch with Uniform B for 20 Iterations:

Transition Matrix A after 20 Iterations:

```
[[0.5 0.5]
 [0.5 0.5]]
```

Emission Matrix B after 20 Iterations:

```
[[0.285 0.19 0.525]
 [0.285 0.19 0.525]]
```

Log-Likelihood after 20 Iterations:

-202.31544020689537

Running Baum-Welch with Uniform B for 100 Iterations:

Transition Matrix A after 100 Iterations:

```
[[0.5 0.5]
 [0.5 0.5]]
```

Emission Matrix B after 100 Iterations:

```
[[0.285 0.19 0.525]
 [0.285 0.19 0.525]]
```

Log-Likelihood after 100 Iterations:

-202.31544020689537

Running Baum-Welch with Uniform B for 20 Iterations:

Transition Matrix A after 20 Iterations:

```
[[0.5 0.5]
 [0.5 0.5]]
```

Emission Matrix B after 20 Iterations:

```
[[0.285 0.19 0.525]
 [0.285 0.19 0.525]]
```

Log-Likelihood after 20 Iterations:

-202.31544020689537

Running Baum-Welch with Uniform B for 100 Iterations:

Transition Matrix A after 100 Iterations:

```
[[0.5 0.5]
 [0.5 0.5]]
```

Emission Matrix B after 100 Iterations:

```
[[0.285 0.19 0.525]
```


[0.285 0.19 0.525]]

Log-Likelihood after 100 Iterations:

-202.31544020689537

The algorithm stabilizes after exactly one iteration. This makes sense as algorithm is unable to preference any of the hidden states as their outcomes are identical. This doesn't even require a constant value across the matrix, just as long as the rows of B are identical.