

# Assignment5\_9628\_jzeiders

December 1, 2024

## 1 John Zeiders (jzeiders) - Assignment 5

```
[ ]: # Import required libraries
import numpy as np
import pandas as pd
from sklearn.metrics import confusion_matrix

# Set random seed for reproducibility
np.random.seed(42)
```

```
[ ]: class PegasosSVM:
    def __init__(self, lambda_param=0.01, n_epochs=20, random_state=42):
        """
        Initialize Pegasos SVM classifier

        Args:
            lambda_param (float): Regularization parameter
            n_epochs (int): Number of epochs for training
            random_state (int): Random seed for reproducibility
        """
        self.lambda_param = lambda_param
        self.n_epochs = n_epochs
        self.random_state = random_state
        self.beta = None

    def _shuffle_data(self, X, y):
        """Shuffle the data while maintaining correspondence between X and y"""
        indices = np.arange(len(y))
        np.random.shuffle(indices)
        return X[indices], y[indices]

    def fit(self, X, y):
        """
        Fit the SVM model using the Pegasos algorithm

        Args:
            X (np.ndarray): Training features of shape (n_samples, n_features)
        """
```

```

        y (np.ndarray): Training labels of shape (n_samples,)
    """
    np.random.seed(self.random_state)

    n_samples, n_features = X.shape

    # Initialize beta (including bias term)
    self.beta = np.zeros(n_features)

    t = 0 # Iteration counter

    # Training loop over epochs
    for epoch in range(self.n_epochs):
        # Shuffle data at the beginning of each epoch
        X_shuffled, y_shuffled = self._shuffle_data(X, y)

        # Loop over each training example
        for i in range(n_samples):
            t += 1
            eta_t = 1 / (self.lambda_param * t) # Learning rate

            # Calculate the decision value
            decision_value = y_shuffled[i] * np.dot(X_shuffled[i], self.
↪beta)

            # Check for margin violation
            if decision_value < 1:
                # Update beta with both gradient and regularization
                self.beta = (1 - eta_t * self.lambda_param) * self.beta + \
                    eta_t * y_shuffled[i] * X_shuffled[i]
            else:
                # Update beta with only regularization
                self.beta = (1 - eta_t * self.lambda_param) * self.beta

            # Projection step
            norm_beta = np.linalg.norm(self.beta)
            threshold = 1 / np.sqrt(self.lambda_param)
            if norm_beta > threshold:
                self.beta = (threshold / norm_beta) * self.beta

        # Print progress
        if (epoch + 1) % 5 == 0 or epoch == 0:
            train_acc = np.mean(self.predict(X) == y)
            print(f"Epoch {epoch + 1}/{self.n_epochs}, Training Accuracy:
↪{train_acc:.4f}")

    def predict(self, X):

```

```

    """
    Predict class labels for samples in X

    Args:
        X (np.ndarray): Features of shape (n_samples, n_features)

    Returns:
        np.ndarray: Predicted class labels (-1 or 1)
    """
    scores = np.dot(X, self.beta)
    return np.where(scores > 0, 1, -1)

def evaluate(self, X, y):
    """
    Evaluate the model and print confusion matrix

    Args:
        X (np.ndarray): Features
        y (np.ndarray): True labels

    Returns:
        float: Classification error rate
    """
    y_pred = self.predict(X)
    conf_matrix = confusion_matrix(y, y_pred, labels=[-1, 1])
    error_rate = np.mean(y_pred != y)

    print("\nConfusion Matrix:")
    print("Predicted")
    print("    -1    1")
    print("-1", conf_matrix[0])
    print(" 1", conf_matrix[1])
    print(f"\nError Rate: {error_rate:.4f}")

    return error_rate

```

```

[ ]: def load_and_preprocess_data(filepath):
    """
    Load and preprocess the data

    Args:
        filepath (str): Path to the CSV file

    Returns:
        tuple: Preprocessed features and labels
    """
    data = pd.read_csv(filepath)

```

```

X = data.iloc[:, :-1].values # All columns except the last

# Add bias term
X = np.hstack((X, np.ones((X.shape[0], 1))))

# Convert labels
y = data.iloc[:, -1].values
y[y == 5] = -1 # Convert 5 to -1
y[y == 6] = 1  # Convert 6 to 1

# Validate labels
unique_labels = np.unique(y)
if set(unique_labels) != {-1, 1}:
    raise ValueError(f"Labels must be -1 and 1, but found {unique_labels}")

return X, y

```

```

[ ]: # URLs to datasets
train_url = 'https://liangfgithub.github.io/Data/coding5_train.csv'
test_url = 'https://liangfgithub.github.io/Data/coding5_test.csv'

# Load and preprocess data
print("Loading and preprocessing data...")
X_train, y_train = load_and_preprocess_data(train_url)
X_test, y_test = load_and_preprocess_data(test_url)

# Initialize and train model
svm = PegasosSVM(lambda_param=0.01, n_epochs=20, random_state=42)

print("\nTraining SVM...")
svm.fit(X_train, y_train)

# Evaluate model
print("\nTraining Set Evaluation:")
train_error = svm.evaluate(X_train, y_train)

print("\nTest Set Evaluation:")
test_error = svm.evaluate(X_test, y_test)

```

Loading and preprocessing data...

Training SVM...

Epoch 1/20, Training Accuracy: 0.7100  
Epoch 5/20, Training Accuracy: 0.9700  
Epoch 10/20, Training Accuracy: 1.0000  
Epoch 15/20, Training Accuracy: 1.0000  
Epoch 20/20, Training Accuracy: 1.0000

Training Set Evaluation:

Confusion Matrix:

Predicted

	-1	1
-1	100	0
1	0	100

Error Rate: 0.0000

Test Set Evaluation:

Confusion Matrix:

Predicted

	-1	1
-1	280	20
1	6	294

Error Rate: 0.0433